



# Flower recognition

---

*Empirical Research 2019/2020*

*Andrea Rafanelli*  
4815168

# Agenda

- 1 DATASET AND DATA PREPROCESSING
- 2 MODEL CREATION
- 3 MODEL IMPROVEMENTS
- 4 WHAT ELSE?
- 5 CONCLUSION

## **Section 1**

# **DATASET AND DATA PREPROCESSING**

# Dataset

GOAL:

*Creation of a model able to classifies flower images into the five categories.*

- 1 4323 images
- 2 Five categories:  
dandelion, rose, sunflower, daisy  
and tulip.
- 3 240x320 pixels






# Data preprocessing (1)

## THREE FILTERS: RGB

```
channels = ['r','g','b']
```

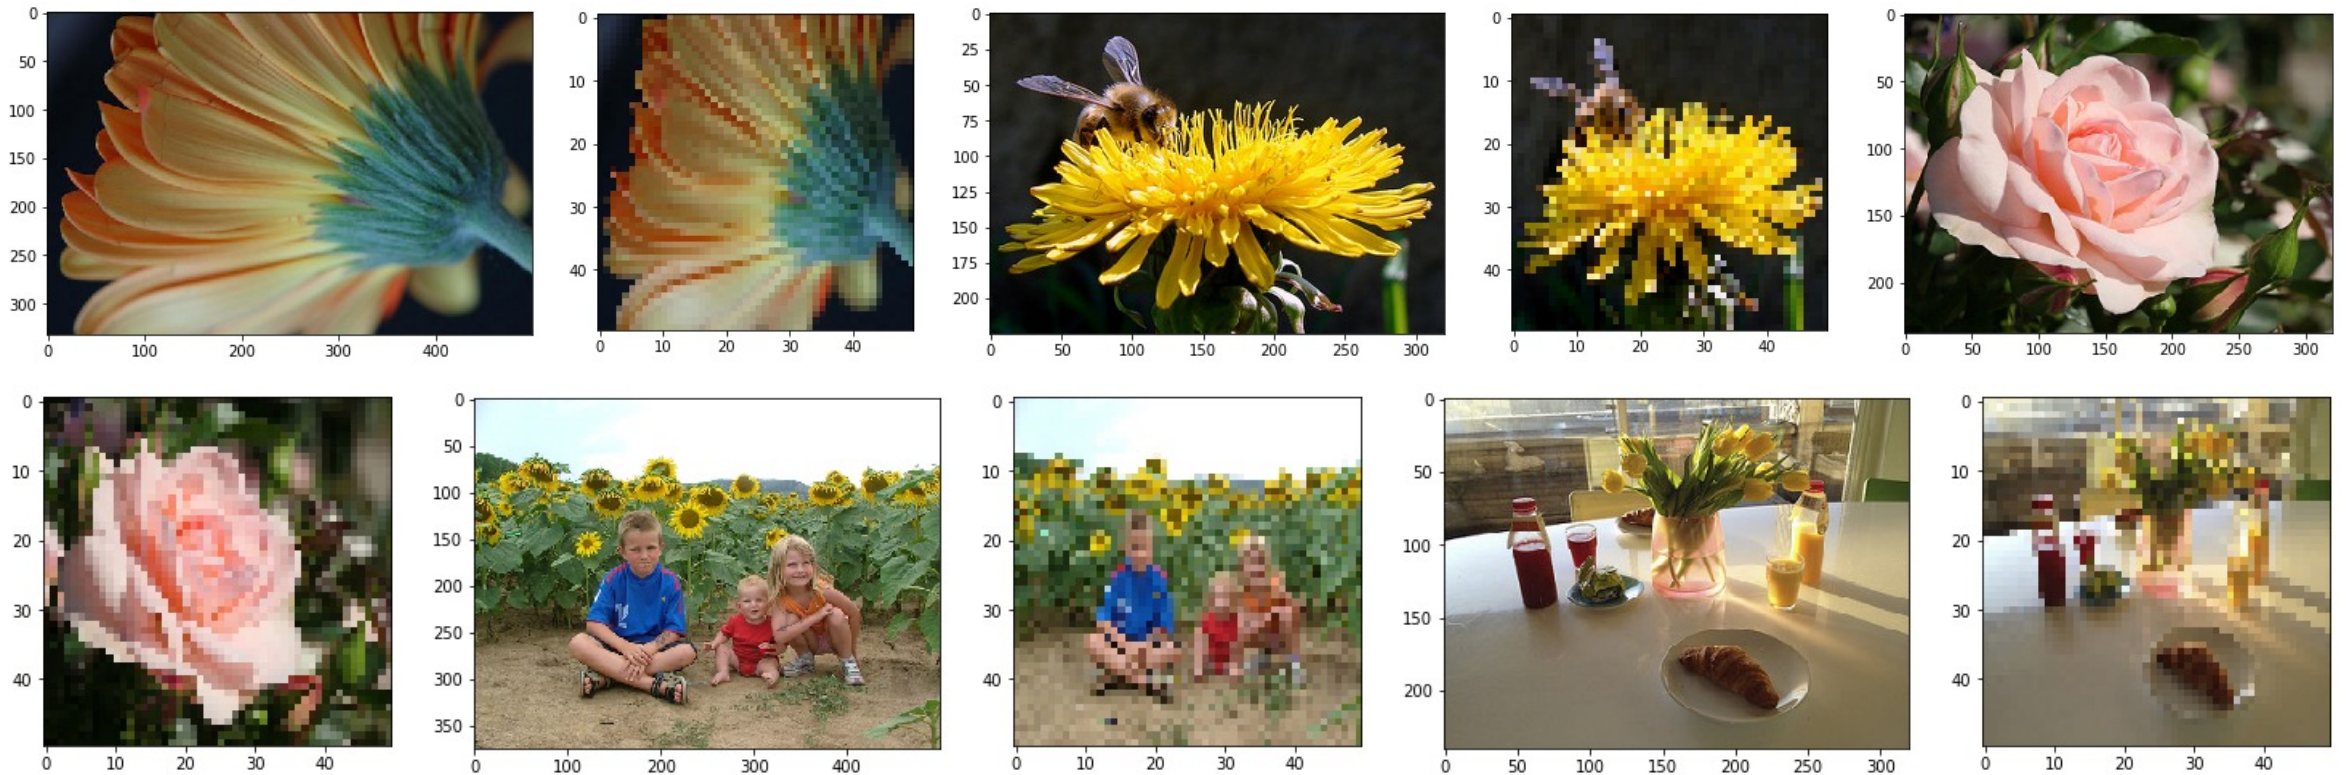
```
def plot_rgb (image):  
    for i, color in enumerate(channels):  
        plt.imshow(image[:, :, i])  
        plt.show()
```



# Data preprocessing (2)


FROM 240x320 TO 90x90

NORMALIZATION OF THE PIXELS




# Data preprocessing (3)

- 3458 images in the training set
- 865 images in the test set



*Are the two balanced?*




## TEST SET

- 18% daisy
- 24% dandelion
- 18% rose
- 17% sunflower
- 23% tulip

## TRAINING SET

- 18% daisy
- 24% dandelion
- 18% rose
- 17% sunflower
- 23% tulip




## **Section 2**

# MODEL CREATION



# Starting with the simplest one

- *Hidden layer activation: **Relu***
- *Weight initialization: '**He initialization**'*  
Because of the non-linearity of the Relu activation function.
- *Output layer activation: **Soft Max***
- *Loss: **Categorical cross entropy***
- *32 as batch size and 10 epochs*
- *3 convolutional layers*
- *Max pooling: extraction of the most important features.*

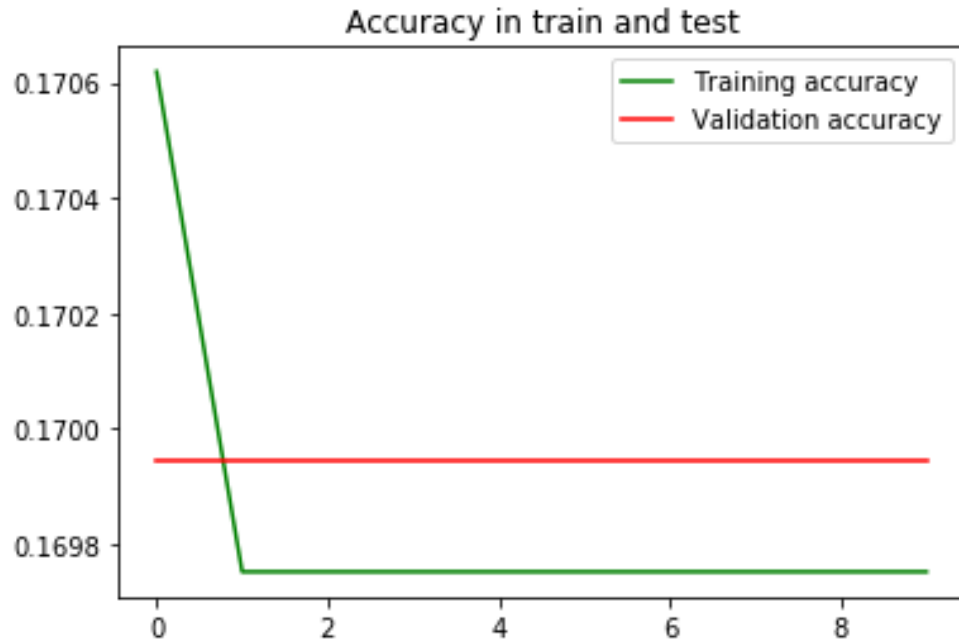


Siddharth Krishna Kumar proves mathematically that for the ReLU activation function, the best weight initialization strategy is to initialize the weights randomly but with this variance:


$$v2=2/N$$

This is exactly the He initialization.

# The ignorance of my model



ADAM






SGD

*Why my model does not learn?*

# From Relu to Leaky Relu

- 1 **Relu** is linear for all positive values, and zero for all negative values.
- 2 **PROBLEM: “dying Relu”**  
The weights and the bias causing the negative preactivations cannot be updated.
- 3 **SOLUTION: Leaky Relu**  
The backward pass is able to alter weights which produce a negative preactivation as the gradient of the activation.



# Which algorithm to use? (1)


## Adaptive Moment Estimation (Adam)

It chooses a direction that would lower the error rate, and continue iterating until the objective function converges to the minimum.


### Learning rate:

- 0.01
- 0.001
- 0.0001
- 0.00001

*L.r: 0.001  
98% accuracy in  
the training set  
vs 68%.*



*L.r: 0.0001  
99% accuracy in  
the training set  
vs 70%.*





# Which algorithm to use? (2)

## Stochastic Gradient Descent (SGD)

It does not perform computation on the whole dataset but only on a small subset or random selection of data examples.

### Learning rate:


- 0.01
- 0.001
- 0.0001
- 0.00001

### Momentum:


To improve the stability, the convergence and the speed of the training.

- 0.5
- 0.9


*L.r: 0.001  
Momentum: 0.5  
53% accuracy in  
the training set  
vs 48%.*



*L.r: 0.0001  
Momentum: 0.9  
60% accuracy in  
the training set  
vs 52%.*




# Which algorithm to use? (3)



**ADAM**

Best model:

Learning rate 0.0001



**SGD**

Best model:

Learning rate 0.0001

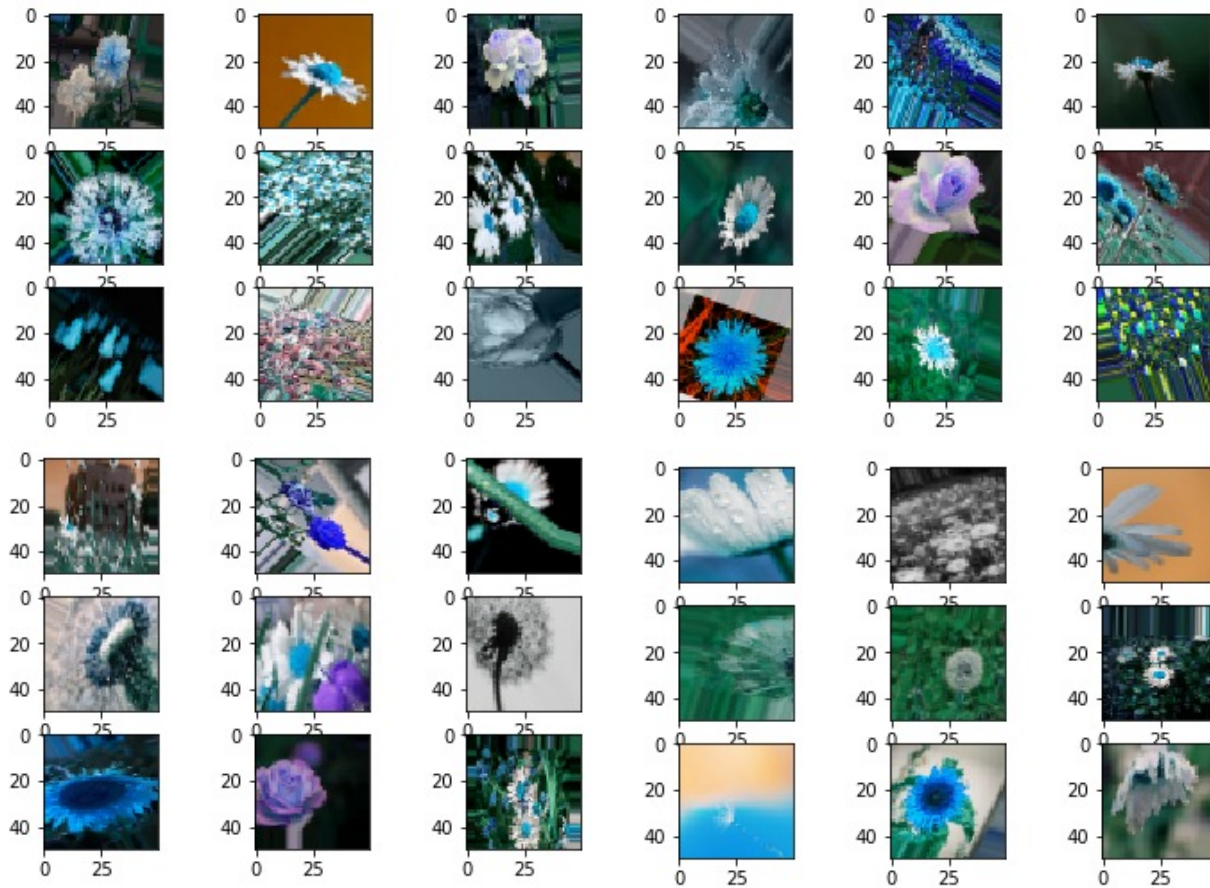
Momentum 0.9

*SGD has less overfitting, thus i decide to use it!*

## Section 3

# MODEL IMPROVEMENTS


# Avoiding overfitting: data augmentation (1)



- Flipping the image horizontally
- Rotation  
Rotation of the image of  $40^\circ$ .
- Width shift  
Shifting the image horizontally by 20%.
- Height shift  
Shifting the image vertically by 20%.
- Zoom  
Zooming the image of 30 %.



# Avoiding overfitting: data augmentation (2)



*69% accuracy in training set and 67% in test set  
No more overfitting but the learning is unstable!*

# Increasing the stability: Batch normalization (1)

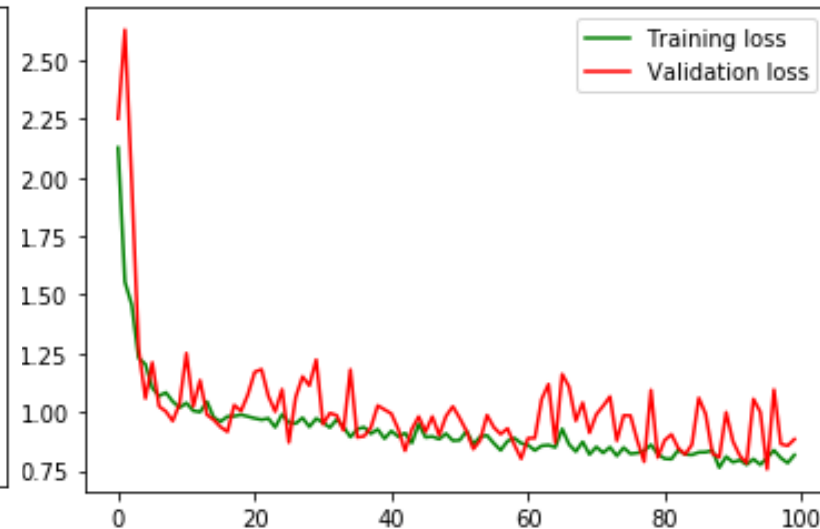
*Not a large improvement:*

*69% accuracy in training set and 68% in test set.*

Accuracy in train and test



Loss in train and test

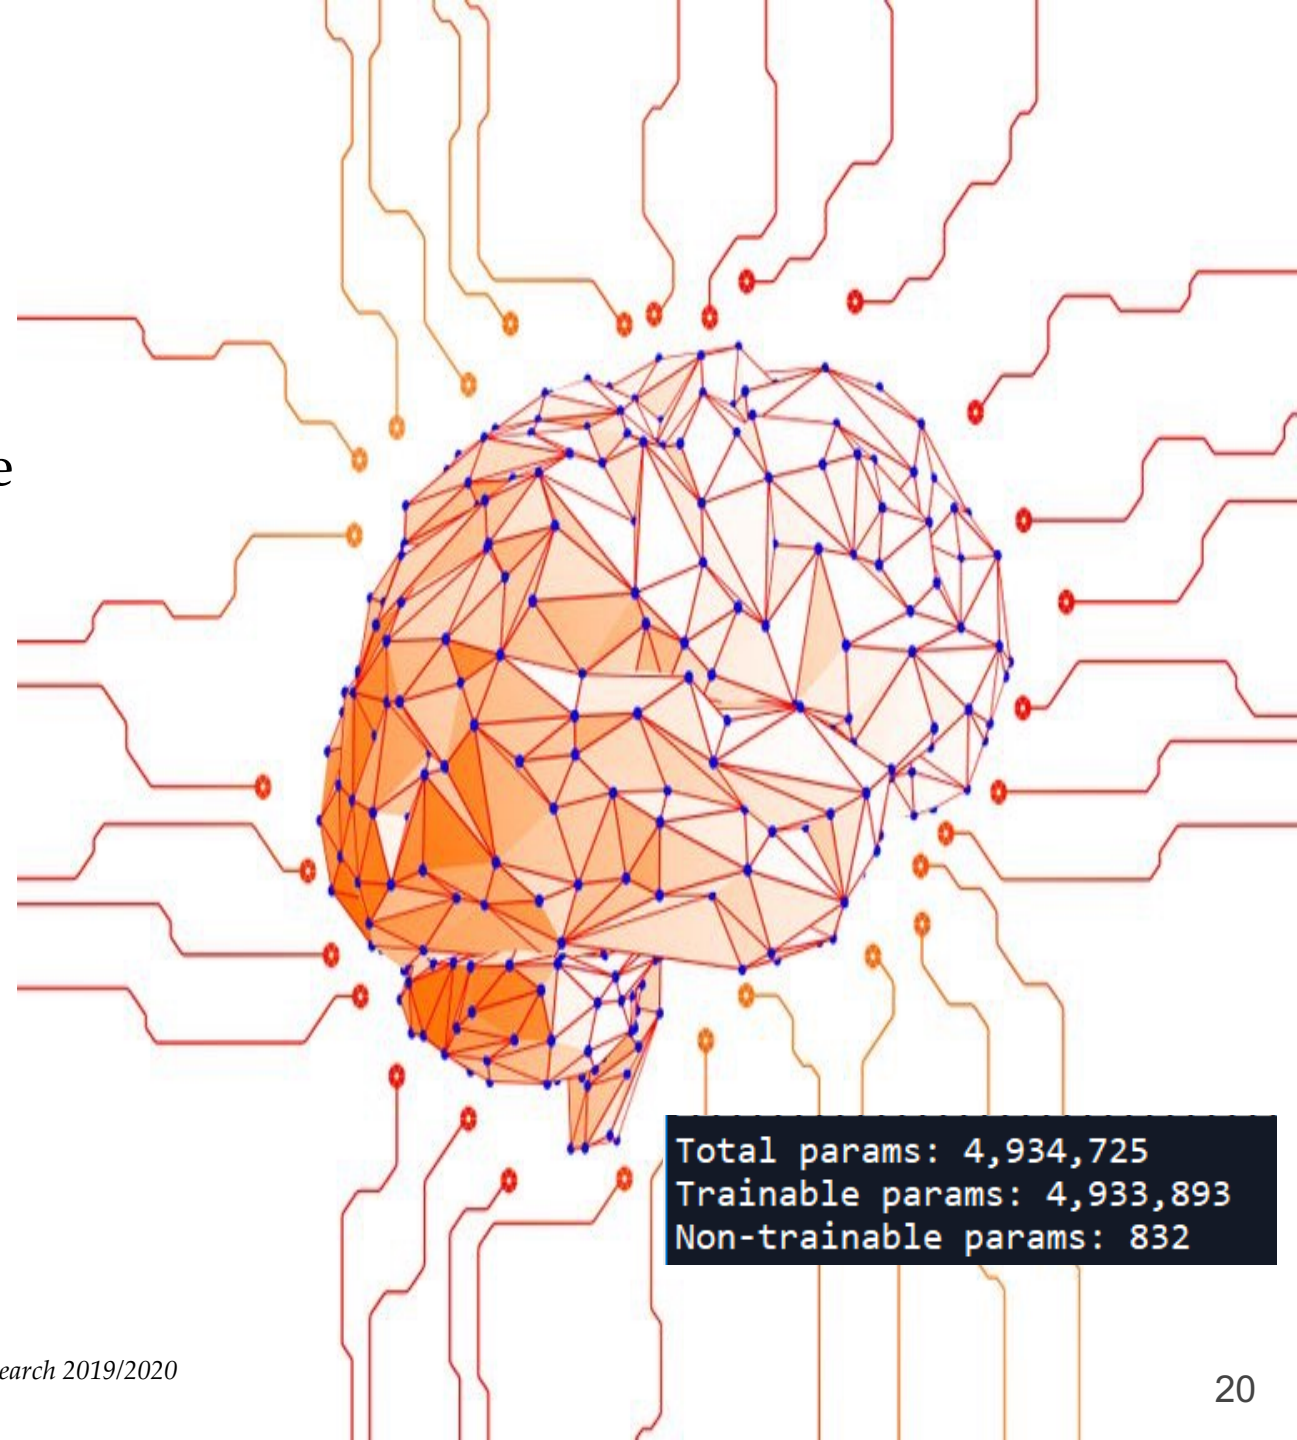


- Improvement of stability and performance
- It normalizes the output of the previous activation layer by subtracting the batch mean and dividing by the batch standard deviation.
- After each convolutional layer

# Final model


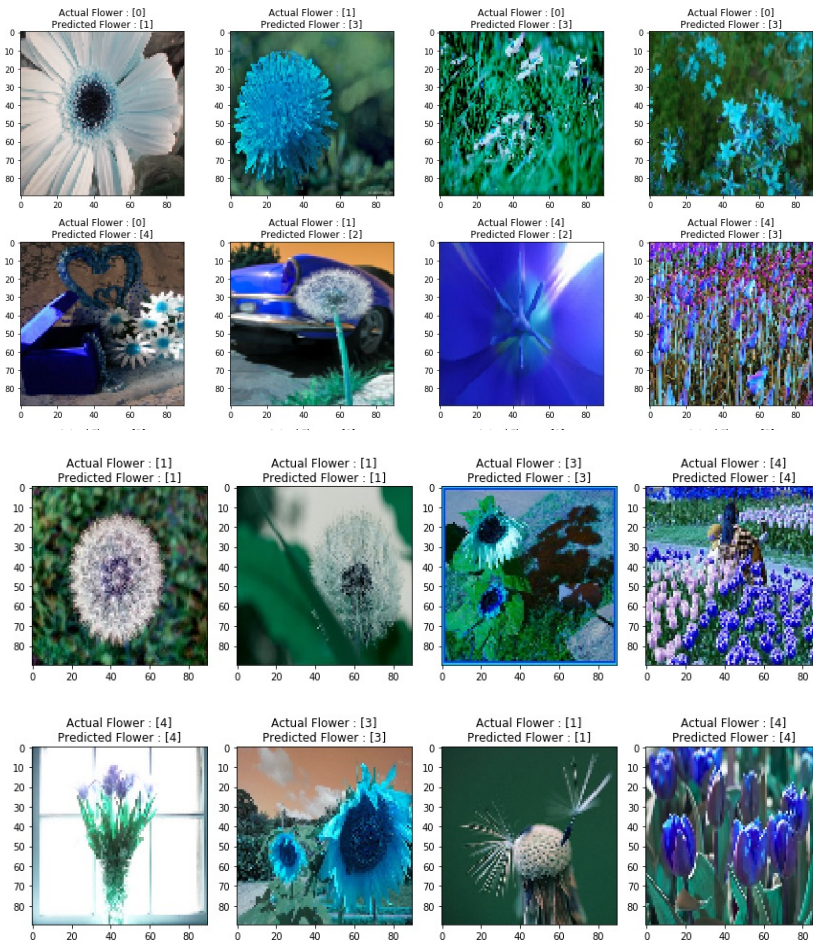
## SUMMARY:

- Optimizer: *SGD* with 0.0001 as learning rate and 0.9 as momentum;
- Activation functions:
  - *Leaky Relu* in the hidden layers
  - *Soft Max* in the output layers
- *Categorical cross-entropy* as loss function;
- 32 as batch size and 100 epochs;
- *Data augmentation* + *Batch normalization*;



# Prediction

*580 correct classification (67%)  
285 uncorrect (33%).*



*Empirical research 2019/2020*



## Section 4

# WHAT ELSE?


A close-up shot of Leonardo DiCaprio in a dark suit, white shirt, and dark tie. He has a serious expression and is looking slightly to his right. In the background, another man with dark hair is partially visible, looking down. The lighting is warm and dramatic, typical of a movie scene. The text "WE NEED TO GO" is overlaid in large, white, bold, sans-serif font with a black outline at the top of the image.

**WE NEED TO GO**

**DEEPER**

# Adding complexity

*Increasing the depth to increase the “levels” of features: higher accuracy*



Vanishing Gradient



Exploding Gradient

MAIN PROBLEM:  
Vanishing/exploding gradient

HOW TO SOLVE IT?


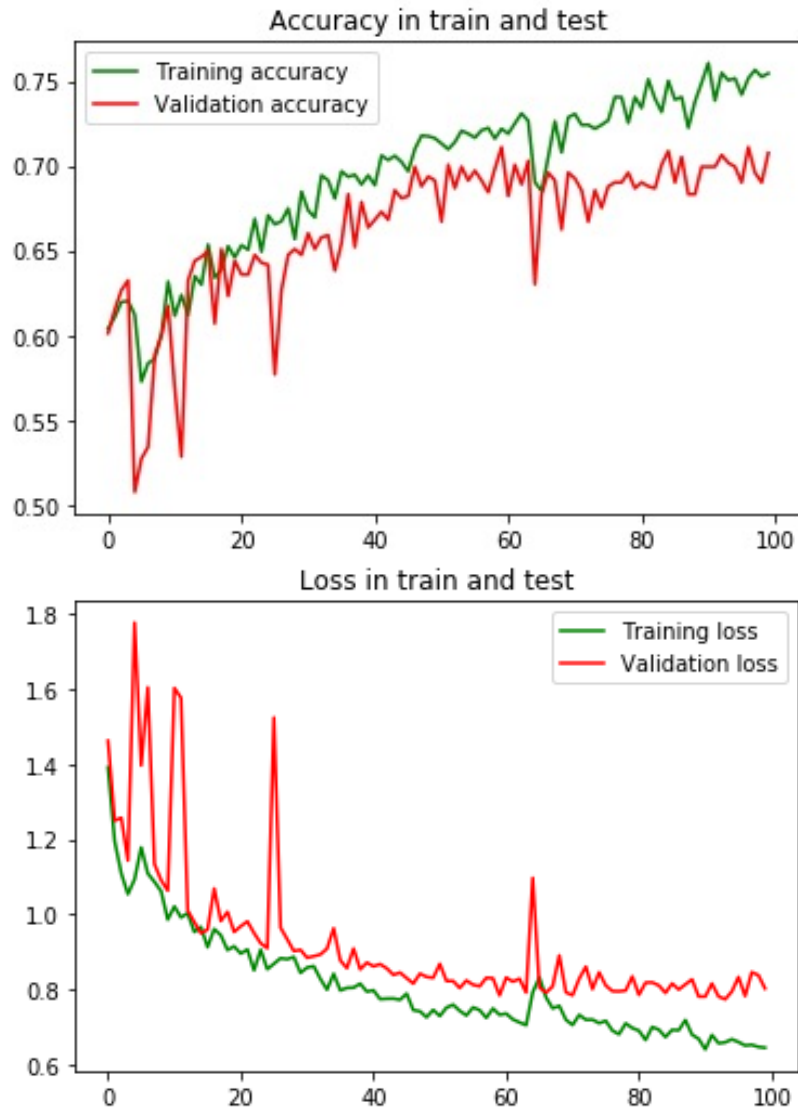


Figure 2. Residual learning: a building block.

# ResNet-50



- It is trained on more than a million images from the ImageNet database.
- He initializer, batch normalization, no dropout.
- Mitigation of vanishing gradient problem because of skip connections.

Skip connections are extra connections between nodes in different layers of a neural network that skip one or more layers.
- It reuses activations from a previous layer until the adjacent layer learns its weights.

*75% accuracy in training set and 69% in test set.*




# Prediction

620 correct classification (72%)  
245 uncorrect (28%).

```
In [67]: RES50.summary()  
Model: "sequential_3"
```

Layer (type)	Output Shape	Param #
resnet50 (Model)	(None, 2, 2, 2048)	23587712
flatten_1 (Flatten)	(None, 8192)	0
batch_normalization_2 (Batch Normalization)	(None, 8192)	32768
dense_2 (Dense)	(None, 2000)	16386000
batch_normalization_3 (Batch Normalization)	(None, 2000)	8000
dense_3 (Dense)	(None, 5)	10005
Total params: 40,024,485		
Trainable params: 39,950,981		
Non-trainable params: 73,504		



*We have not a strong improvement. Why?*

- Too much parameters to estimate and few images.
- Because of the large number of weights that are not updated during the train.

## **Section 5**

# CONCLUSION

# Conclusion

- Not always deeper is better : Resnet-50 is very complex and it works well for large dataset:  
In this case is better to work with «customized» model.
- Try different combination and structures: average pooling, Adam optimizer, more/less convolutional layers etc.
- Use of *cv* to find the best parameters ( learning rate, batch size, number of epochs, optimizer..)
- Use of *GPU*: maintain the original size.

# References

- *Chollet François (2017), «Deep learning with Python»*
- *Ioffe Sergey, Szegedy Christian (2015), "Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift"*
- *Kaiming He et al. (2015), Deep Residual Learning for Image Recognition»*
- *Kaiming He et al. (2015), "Delving Deep into Rectifiers: Surpassing Human-Level Performance on ImageNet Classification"*
- *Krizhevsky Alex, Ilya Sutskever and Geoffrey E. Hinton (2012), "Imagenet classification with deep convolutional neural networks"*